

Analysis of Linux Driver and Device matching process



garden of idyllic poets

Modified at 2022-05-29 23:33:02

1035

Favorite 9

copyright

Category column: linux embedded development

Article tags: linux

Device and driver matching process analysis



linux embedded dev... The column includes this content

Analysis of Linux Driver and Device matching process

1. Bus registration matching process
 - 1.1 struct platform_driver
 - 1.2 struct device_driver
 - 1.3 The process of registering the PCIe bus controller driver to the platform
 - 1.3.1 rockchip_pcie_driver
 - 1.3.2 module_platform_driver
 - 1.3.3 platform_driver_register
 - 1.3.4 __platform_driver_register and platform_bus_type
 - 1.3.5 driver_register
 - 1.3.6 bus_add_driver
 - 1.3.7 driver_attach
 - 1.3.8 bus_for_each_dev
 - 1.3.9 __driver_attach
 - 1.3.10 driver_match_device
 - 1.3.11 platform_match
 - 1.3.12 driver_probe_device
 - 2.1.12 __driver_probe_device
 - 1.3.13 really_probe
 - 1.3.14 call_driver_probe
 - 1.3.15 platform_drv_probe
 - 1.3.16 rockchip_pcie_probe
2. Device driver registration matching process:
 - 2.1 pci_register_driver
 - 2.1.1 nvme_init
 - 2.1.2 pci_register_driver
 - 2.1.3 __pci_register_driver
 - 2.1.4 driver_register
 - 2.1.5 bus_add_driver
 - 2.1.6 driver_attach
 - 2.1.7 bus_for_each_dev
 - 2.1.8 __driver_attach
 - 2.1.9 driver_match_device
 - 2.1.10 pci_bus_match
 - 2.1.11 pci_match_device
 - 2.1.12 pci_match_one_device
 - 2.1.13 driver_probe_device
 - 2.1.12 __driver_probe_device
 - 2.1.14 really_probe
 - 2.1.15 call_driver_probe
 - 2.1.16 pci_bus_type
 - 2.1.17 pci_device_probe
 - 2.1.18 __pci_device_probe
 - 2.1.19 pci_call_probe
 - 2.1.20 local_pci_probe
 - 2.1.21 nvme_probe

Take the linux-5.14.10 kernel as an example to analyze the bus registration process:

it is based on the processing flow analysis of the device tree, but the mode of bus/dev/drv is only slightly different in the matching process of the match function, and the other processes are the same.

1. Bus registration matching process

The bus registration process takes the PCIe controller registration process of rk3399 as an example for analysis.

1.1 struct platform_driver

code path: `include/linux/platform_device.h`

```

1      206 struct platform_driver {
2      207     int (*probe)(struct platform_device *);
3      208     int (*remove)(struct platform_device *);
4      209     void (*shutdown)(struct platform_device *);
5      210     int (*suspend)(struct platform_device *, pm_message_t state);
6      211     int (*resume)(struct platform_device *);
7      212     struct device_driver driver;
8      213     const struct platform_device_id *id_table;
9      214     bool prevent_deferred_probe;
10     215 };

```

1.2 struct device_driver

code path: `include/linux/device/device.h`

```

1      50 /**
2      51  * struct device_driver - The basic device driver structure
3      52  * @name: Name of the device driver.
4      53  * @bus: The bus which the device of this driver belongs to.
5      54  * @owner: The module owner.
6      55  * @mod_name: Used for built-in modules.
7      56  * @suppress_bind_attrs: Disables bind/unbind via sysfs.
8      57  * @probe_type: Type of the probe (synchronous or asynchronous) to use.
9      58  * @of_match_table: The open firmware table.
10     59  * @acpi_match_table: The ACPI match table.
11     60  * @probe: Called to query the existence of a specific device,
12     61  *         whether this driver can work with it, and bind the driver
13     62  *         to a specific device.
14     63  * @sync_state: Called to sync device state to software state after all the
15     64  *         state tracking consumers linked to this device (present at
16     65  *         the time of late_initcall) have successfully bound to a
17     66  *         driver. If the device has no consumers, this function will
18     67  *         be called at late_initcall_sync level. If the device has
19     68  *         consumers that are never bound to a driver, this function
20     69  *         will never get called until they do.
21     70  * @remove: Called when the device is removed from the system to
22     71  *         unbind a device from this driver.
23     72  * @shutdown: Called at shut-down time to quiesce the device.
24     73  * @suspend: Called to put the device to sleep mode. Usually to a
25     74  *         low power state.
26     75  * @resume: Called to bring a device from sleep mode.
27     76  * @groups: Default attributes that get created by the driver core
28     77  *         automatically.
29     78  * @dev_groups: Additional attributes attached to device instance once
30     79  *         it is bound to the driver.
31     80  * @pm: Power management operations of the device which matched
32     81  *         this driver.
33     82  * @coredump: Called when sysfs entry is written to. The device driver
34     83  *         is expected to call the dev_coredump API resulting in a
35     84  *         uevent.
36     85  * @p: Driver core's private data, no one other than the driver
37     86  *         core can touch this.
38     87  *
39     88  * The device driver-model tracks all of the drivers known to the system.
40     89  * The main reason for this tracking is to enable the driver core to match
41     90  * up drivers with new devices. Once drivers are known objects within the
42     91  * system, however, a number of other things become possible. Device drivers
43     92  * can export information and configuration variables that are independent
44     93  * of any specific device.
45     94  */
46     95 struct device_driver {
47     96     const char *name;
48     97     struct bus_type *bus;
49     98
50     99     struct module *owner;
51    100     const char *mod_name; /* used for built-in modules */
52    101
53    102     bool suppress_bind_attrs; /* disables bind/unbind via sysfs */
54    103     enum probe_type probe_type;
55    104
56    105     const struct of_device_id *of_match_table;
57    106     const struct acpi_device_id *acpi_match_table;

```

```

47 | 107
48 | 108     int (*probe) (struct device *dev);
49 | 109     void (*sync_state)(struct device *dev);
50 | 110     int (*remove) (struct device *dev);
51 | 111     void (*shutdown) (struct device *dev);
52 | 112     int (*suspend) (struct device *dev, pm_message_t state);
53 | 113     int (*resume) (struct device *dev);
54 | 114     const struct attribute_group **groups;
55 | 115     const struct attribute_group **dev_groups;
56 | 116
57 | 117     const struct dev_pm_ops *pm;
58 | 118     void (*coredump) (struct device *dev);
59 | 119
60 | 120     struct driver_private *p;
61 | 121 };
62 |

```

1.3 The process of registering the PCIe bus controller driver to the platform

```

1 | |- module_platform_driver
2 |   |- platform_driver_register
3 |     |- __platform_driver_register
4 |       |- driver_register
5 |         |- bus_add_driver
6 |           |- driver_attach
7 |             |- bus_for_each_dev
8 |               |- __driver_attach
9 |                 |- driver_match_device
10 |                   |- platform_match
11 |                     |- driver_probe_device
12 |                       |- really_probe
13 |                         |- call_driver_probe
14 |                           |- platform_drv_probe
15 |                             |- rockchip_pcie_probe

```

1.3.1 rockchip_pcie_driver

code path: `drivers/pci/controller/pcie-rockchip-host.c`

```

1 | 1041 static const struct dev_pm_ops rockchip_pcie_pm_ops = {
2 | 1042     SET_NOIRQ_SYSTEM_SLEEP_PM_OPS(rockchip_pcie_suspend_noirq,
3 | 1043                                     rockchip_pcie_resume_noirq)
4 | 1044 };
5 | 1045
6 | 1046 static const struct of_device_id rockchip_pcie_of_match[] = {
7 | 1047     { .compatible = "rockchip,rk3399-pcie", },
8 | 1048     {}
9 | 1049 };
10 | 1050 MODULE_DEVICE_TABLE(of, rockchip_pcie_of_match);
11 | 1051
12 | 1052 static struct platform_driver rockchip_pcie_driver = {
13 | 1053     .driver = {
14 | 1054         .name = "rockchip-pcie",
15 | 1055         .of_match_table = rockchip_pcie_of_match,
16 | 1056         .pm = &rockchip_pcie_pm_ops,
17 | 1057     },
18 | 1058     .probe = rockchip_pcie_probe,
19 | 1059     .remove = rockchip_pcie_remove,
20 | 1060 };
21 | 1061 module_platform_driver(rockchip_pcie_driver);

```

1.3.2 module_platform_driver

code path: `include/linux/platform_device.h`

```

1 |
2 | 248 /* module_platform_driver() - Helper macro for drivers that don't do
   | 249 * anything special in module init/exit. This eliminates a lot of

```

```

3 | 250 * boilerplate. Each module may only use this macro once, and
4 | 251 * calling it replaces module_init() and module_exit()
5 | 252 */
6 | 253 #define module_platform_driver(__platform_driver) \
7 | 254     module_driver(__platform_driver, platform_driver_register, \
8 | 255     platform_driver_unregister)

```

1.3.3 platform_driver_register

code path: `include/linux/platform_device.h`

```

1 | 220 /*
2 | 221 * use a macro to avoid include chaining to get THIS_MODULE
3 | 222 */
4 | 223 #define platform_driver_register(drv) \
5 | 224     __platform_driver_register(drv, THIS_MODULE)

```

1.3.4 __platform_driver_register and platform_bus_type

code path: `drivers/base/platform.c`

```

1 | 863 /**
2 | 864 * __platform_driver_register - register a driver for platform-level devices
3 | 865 * @drv: platform driver structure
4 | 866 * @owner: owning module/driver
5 | 867 */
6 | 868 int __platform_driver_register(struct platform_driver *drv,
7 | 869     struct module *owner)
8 | 870 {
9 | 871     drv->driver.owner = owner;
10 | 872     drv->driver.bus = &platform_bus_type;
11 | 873
12 | 874     return driver_register(&drv->driver);
13 | 875 }
14 | 876 EXPORT_SYMBOL_GPL(__platform_driver_register);

```

code path: `drivers/base/platform.c`

```

1 | 1492 struct bus_type platform_bus_type = {
2 | 1493     .name      = "platform",
3 | 1494     .dev_groups = platform_dev_groups,
4 | 1495     .match      = platform_match,
5 | 1496     .uevent     = platform_uevent,
6 | 1497     .probe      = platform_probe,
7 | 1498     .remove     = platform_remove,
8 | 1499     .shutdown   = platform_shutdown,
9 | 1500     .dma_configure = platform_dma_configure,
10 | 1501     .pm         = &platform_dev_pm_ops,
11 | 1502 };
12 | 1503 EXPORT_SYMBOL_GPL(platform_bus_type);

```

1.3.5 driver_register

code path: `drivers/base/driver.c`

```

1 | 139 /**
2 | 140 * driver_register - register driver with bus
3 | 141 * @drv: driver to register
4 | 142 *
5 | 143 * We pass off most of the work to the bus_add_driver() call,
6 | 144 * since most of the things we have to do deal with the bus
7 | 145 * structures.
8 | 146 */
9 | 147 int driver_register(struct device_driver *drv)
10 | 148 {
11 | 149     int ret;
12 | 150     struct device_driver *other;

```

```

12 151
13 152     if (!drv->bus->p) {
14 153         pr_err("Driver '%s' was unable to register with bus_type '%s' because the bus was not initialized.\n",
15 154             drv->name, drv->bus->name);
16 155         return -EINVAL;
17 156     }
18 157
19 158     if ((drv->bus->probe && drv->probe) ||
20 159         (drv->bus->remove && drv->remove) ||
21 160         (drv->bus->shutdown && drv->shutdown))
22 161         pr_warn("Driver '%s' needs updating - please use "
23 162             "bus_type methods\n", drv->name);
24 163
25 164     other = driver_find(drv->name, drv->bus);
26 165     if (other) {
27 166         pr_err("Error: Driver '%s' is already registered, "
28 167             "aborting...\n", drv->name);
29 168         return -EBUSY;
30 169     }
31 170
32 171     ret = bus_add_driver(drv);
33 172     if (ret)
34 173         return ret;
35 174     ret = driver_add_groups(drv, drv->groups);
36 175     if (ret) {
37 176         bus_remove_driver(drv);
38 177         return ret;
39 178     }
40 179     kobject_uevent(&drv->p->kobj, KOBJ_ADD);
41 180
42 181     return ret;
43 182 }
44 183 EXPORT_SYMBOL_GPL(driver_register);
45

```

1.3.6 bus_add_driver

code path: [drivers/base/bus.c](#)

```

1
2 586 /**
3 587  * bus_add_driver - Add a driver to the bus.
4 588  * @drv: driver.
5 589  */
6 590 int bus_add_driver(struct device_driver *drv)
7 591 {
8 592     struct bus_type *bus;
9 593     struct driver_private *priv;
10 594     int error = 0;
11 595
12 596     bus = bus_get(drv->bus);
13 597     if (!bus)
14 598         return -EINVAL;
15 599
16 600     pr_debug("bus: '%s': add driver %s\n", bus->name, drv->name);
17 601
18 602     priv = kzalloc(sizeof(*priv), GFP_KERNEL);
19 603     if (!priv) {
20 604         error = -ENOMEM;
21 605         goto out_put_bus;
22 606     }
23 607     klist_init(&priv->klist_devices, NULL, NULL);
24 608     priv->driver = drv;
25 609     drv->p = priv;
26 610     priv->kobj.kset = bus->p->drivers_kset;
27 611     error = kobject_init_and_add(&priv->kobj, &driver_ktype, NULL,
28 612         "%s", drv->name);
29 613     if (error)
30 614         goto out_unregister;
31 615
32 616     klist_add_tail(&priv->knode_bus, &bus->p->klist_drivers);

```

```

33     617     if (drv->bus->p->drivers_autoprobe) {
34     618         error = driver_attach(drv);
35     619         if (error)
36     620             goto out_unregister;
37     621     }
38     622     module_add_driver(drv->owner, drv);
39     623
40     624     error = driver_create_file(drv, &driver_attr_uevent);
41     625     if (error) {
42     626         printk(KERN_ERR "%s: uevent attr (%s) failed\n",
43     627             __func__, drv->name);
44     628     }
45     629     error = driver_add_groups(drv, bus->drv_groups);
46     630     if (error) {
47     631         /* How the hell do we get out of this pickle? Give up */
48     632         printk(KERN_ERR "%s: driver_add_groups(%s) failed\n",
49     633             __func__, drv->name);
50     634     }
51     635
52     636     if (!drv->suppress_bind_attrs) {
53     637         error = add_bind_files(drv);
54     638         if (error) {
55     639             /* Ditto */
56     640             printk(KERN_ERR "%s: add_bind_files(%s) failed\n",
57     641                 __func__, drv->name);
58     642         }
59     643     }
60     644
61     645     return 0;
62     646
63     647 out_unregister:
64     648     kobject_put(&priv->kobj);
65     649     /* drv->p is freed in driver_release() */
66     650     drv->p = NULL;
67     651 out_put_bus:
68     652     bus_put(bus);
69     653     return error;
654 }

```

1.3.7 driver_attach

code path: `drivers/base/dd.c`

```

1     1146 /**
2     1147  * driver_attach - try to bind driver to devices.
3     1148  * @drv: driver.
4     1149  *
5     1150  * Walk the list of devices that the bus has on it and try to
6     1151  * match the driver with each one. If driver_probe_device()
7     1152  * returns 0 and the @dev->driver is set, we've found a
8     1153  * compatible pair.
9     1154  */
10    1155 int driver_attach(struct device_driver *drv)
11    1156 {
12    1157     return bus_for_each_dev(drv->bus, NULL, drv, __driver_attach);
13    1158 }
14    1159 EXPORT_SYMBOL_GPL(driver_attach);

```

1.3.8 bus_for_each_dev

code path: `drivers/base/bus.c`

```

1     269 /**
2     270  * bus_for_each_dev - device iterator.
3     271  * @bus: bus type.
4     272  * @start: device to start iterating from.
5     273  * @data: data for the callback.
6     274  * @fn: function to be called for each device.
7     275  *
8     276  * Iterate over @bus's list of devices, and call @fn for each,
9     277  * passing it @data. If @start is not NULL, we use that device to

```

```

9      278 * begin iterating from.
10     279 *
11     280 * We check the return of @fn each time. If it returns anything
12     281 * other than 0, we break out and return that value.
13     282 *
14     283 * NOTE: The device that returns a non-zero value is not retained
15     284 * in any way, nor is its refcount incremented. If the caller needs
16     285 * to retain this data, it should do so, and increment the reference
17     286 * count in the supplied callback.
18     287 */
19     288 int bus_for_each_dev(struct bus_type *bus, struct device *start,
20     289                     void *data, int (*fn)(struct device *, void *))
21 {
22     291     struct klist_iter i;
23     292     struct device *dev;
24     293     int error = 0;
25
26     294
27     295     if (!bus || !bus->p)
28     296         return -EINVAL;
29
30     297
31     298     klist_iter_init_node(&bus->p->klist_devices, &i,
32     299                         (start ? &start->p->knode_bus : NULL));
33
34     300     while (!error && (dev = next_device(&i)))
35     301         error = fn(dev, data);
36
37     302     klist_iter_exit(&i);
38     303     return error;
39 }
40
41     304 }
42
43     305 EXPORT_SYMBOL_GPL(bus_for_each_dev);

```

1.3.9 __driver_attach

- driver_match_device first determines whether the driver and the bus driver match
- driver_probe_device calls the corresponding bus driver probe function to do the corresponding initialization processing
code path: `drivers/base/dd.c`

```

1
2     1092 static int __driver_attach(struct device *dev, void *data)
3     1093 {
4     1094     struct device_driver *drv = data;
5     1095     int ret;
6
7     1096
8     1097     /*
9     1098      * Lock device and try to bind to it. We drop the error
10    1099      * here and always return 0, because we need to keep trying
11    1100      * to bind to devices and some drivers will return an error
12    1101      * simply if it didn't support the device.
13    1102      *
14    1103      * driver_probe_device() will spit a warning if there
15    1104      * is an error.
16    1105      */
17    1106
18    1107     ret = driver_match_device(drv, dev);
19    1108     if (ret == 0) {
20    1109         /* no match */
21    1110         return 0;
22    1111     } else if (ret == -EPROBE_DEFER) {
23    1112         dev_dbg(dev, "Device match requests probe deferral\n");
24    1113         dev->can_match = true;
25    1114         driver_deferred_probe_add(dev);
26    1115     } else if (ret < 0) {
27    1116         dev_dbg(dev, "Bus failed to match device: %d\n", ret);
28    1117         return ret;
29    1118     } /* ret > 0 means positive match */
30
31    1119
32    1120     if (driver_allows_async_probing(drv)) {
33    1121         /*
34    1122          * Instead of probing the device synchronously we will
35    1123          * probe it asynchronously to allow for more parallelism.
36    1124          *
37    1125          * We only take the device lock here in order to guarantee
38    1126          * that the dev->driver and async_driver fields are protected
39    1127          */
40    1128         dev_dbg(dev, "probing driver %s asynchronously\n", drv->name);
41    1129         device_lock(dev);
42    1130         if (!dev->driver) {

```

```

38 | 1131         get_device(dev);
39 | 1132         dev->p->async_driver = drv;
40 | 1133         async_schedule_dev(__driver_attach_async_helper, dev);
41 | 1134     }
42 | 1135     device_unlock(dev);
43 | 1136     return 0;
44 | 1137 }
45 | 1138
46 | 1139     __device_driver_lock(dev, dev->parent);
47 | 1140     driver_probe_device(drv, dev);
48 | 1141     __device_driver_unlock(dev, dev->parent);
49 | 1142
50 | 1143     return 0;
51 | 1144 }
52 |

```

1.3.10 driver_match_device

For `drv->bus->match(dev, drv)`, the match function corresponds to the `platform_match` function in the previously registered `platform_bus_type` code path: `drivers/base/base.h`

```

1 | 144 static inline int driver_match_device(struct device_driver *drv,
2 | 145         struct device *dev)
3 | 146 {
4 | 147     return drv->bus->match ? drv->bus->match(dev, drv) : 1;
5 | 148

```

1.3.11 platform_match

Determine whether the platform driver and platform device match according to the following process, if one succeeds, that is, the match is successful

- compare `platform_dev.driver_override` with `platform_driver.drv->name`
- Compare `platform_dev.dev.of_node` the `compatible` properties of and `platform_driver.drv->of_match_table`
- compare `platform_dev.name` with `platform_driver.id_table`
- compare `platform_dev.name` with `platform_driver.drv->name`

code path: `drivers/base/platform.c`

```

1 |
2 | 1345 /**
3 | 1346  * platform_match - bind platform device to platform driver.
4 | 1347  * @dev: device.
5 | 1348  * @drv: driver.
6 | 1349  *
7 | 1350  * Platform device IDs are assumed to be encoded like this:
8 | 1351  * "<name><instance>", where <name> is a short description of the type of
9 | 1352  * device, like "pci" or "floppy", and <instance> is the enumerated
10 | 1353  * instance of the device, like '0' or '42'. Driver IDs are simply
11 | 1354  * "<name>". So, extract the <name> from the platform_device structure,
12 | 1355  * and compare it against the name of the driver. Return whether they match
13 | 1356  * or not.
14 | 1357  */
15 | 1358 static int platform_match(struct device *dev, struct device_driver *drv)
16 | 1359 {
17 | 1360     struct platform_device *pdev = to_platform_device(dev);
18 | 1361     struct platform_driver *pdrv = to_platform_driver(drv);
19 | 1362
20 | 1363     /* When driver_override is set, only bind to the matching driver */
21 | 1364     if (pdev->driver_override)
22 | 1365         return !strcmp(pdev->driver_override, drv->name);
23 | 1366
24 | 1367     /* Attempt an OF style match first */
25 | 1368     if (of_driver_match_device(dev, drv))
26 | 1369         return 1;
27 | 1370
28 | 1371     /* Then try ACPI style match */
29 | 1372     if (acpi_driver_match_device(dev, drv))
30 | 1373         return 1;
31 | 1374
32 | 1375     /* Then try to match against the id table */

```



```

30 | 1376     if (pdrv->id_table)
31 | 1377         return platform_match_id(pdrv->id_table, pdev) != NULL;
32 | 1378
33 | 1379     /* fall-back to driver name match */
34 | 1380     return (strcmp(pdev->name, drv->name) == 0);
35 | 1381 }
36 |
37 |

```

1.3.12 driver_probe_device

code path: [drivers/base/dd.c](#)

```

1 | 761 /**
2 | 762  * driver_probe_device - attempt to bind device & driver together
3 | 763  * @drv: driver to bind a device to
4 | 764  * @dev: device to try to bind to the driver
5 | 765  *
6 | 766  * This function returns -ENODEV if the device is not registered, -EBUSY if it
7 | 767  * already has a driver, 0 if the device is bound successfully and a positive
8 | 768  * (inverted) error code for failures from the ->probe method.
9 | 769  *
10 | 770  * This function must be called with @dev lock held. When called for a
11 | 771  * USB interface, @dev->parent lock must be held as well.
12 | 772  *
13 | 773  * If the device has a parent, runtime-resume the parent before driver probing.
14 | 774  */
15 | 775 static int driver_probe_device(struct device_driver *drv, struct device *dev)
16 | 776 {
17 | 777     int trigger_count = atomic_read(&deferred_trigger_count);
18 | 778     int ret;
19 | 779
20 | 780     atomic_inc(&probe_count);
21 | 781     ret = __driver_probe_device(drv, dev);
22 | 782     if (ret == -EPROBE_DEFER || ret == EPROBE_DEFER) {
23 | 783         driver_deferred_probe_add(dev);
24 | 784
25 | 785         /*
26 | 786          * Did a trigger occur while probing? Need to re-trigger if yes
27 | 787          */
28 | 788         if (trigger_count != atomic_read(&deferred_trigger_count) &&
29 | 789             !defer_all_probes)
30 | 790             driver_deferred_probe_trigger();
31 | 791     }
32 | 792     atomic_dec(&probe_count);
33 | 793     wake_up_all(&probe_waitqueue);
34 | 794     return ret;
35 | 795 }
36 |
37 |

```

2.1.12 __driver_probe_device

code path: [drivers/base/dd.c](#)

```

1 |
2 | 730 static int __driver_probe_device(struct device_driver *drv, struct device *dev)
3 | 731 {
4 | 732     int ret = 0;
5 | 733
6 | 734     if (dev->p->dead || !device_is_registered(dev))
7 | 735         return -ENODEV;
8 | 736     if (dev->driver)
9 | 737         return -EBUSY;
10 | 738
11 | 739     dev->can_match = true;
12 | 740     pr_debug("bus: '%s': %s: matched device %s with driver %s\n",
13 | 741             drv->bus->name, __func__, dev_name(dev), drv->name);
14 | 742
15 | 743     pm_runtime_get_suppliers(dev);
16 | 744     if (dev->parent)
17 | 745         pm_runtime_get_sync(dev->parent);
18 | 746
19 | 747     pm_runtime_barrier(dev);
20 | 748     if (initcall_debug)

```

```

twen: 749         ret = really_probe_debug(dev, drv);
twen: 750     else
twen: 751         ret = really_probe(dev, drv);
twen: 752     pm_request_idle(dev);
25: 753
26: 754     if (dev->parent)
27: 755         pm_runtime_put(dev->parent);
28: 756
29: 757     pm_runtime_put_suppliers(dev);
30: 758     return ret;
      759 }

```

1.3.13 really_probe

code path: [drivers/base/dd.c](#)

```

2
3     541 static int really_probe(struct device *dev, struct device_driver *drv)
4     542 {
5     543     bool test_remove = IS_ENABLED(CONFIG_DEBUG_TEST_DRIVER_REMOVE) &&
6     544         !drv->suppress_bind_attrs;
7     545     int ret;
8     546
9     547     if (defer_all_probes) {
10    548         /*
11    549          * Value of defer_all_probes can be set only by
12    550          * device_block_probing() which, in turn, will call
13    551          * wait_for_device_probe() right after that to avoid any races.
14    552          */
15    553         dev_dbg(dev, "Driver %s force probe deferral\n", drv->name);
16    554         return -EPROBE_DEFER;
17    555     }
18    556
19    557     ret = device_links_check_suppliers(dev);
20    558     if (ret)
21    559         return ret;
22    560
23    561     pr_debug("bus: '%s': %s: probing driver %s with device %s\n",
24    562         drv->bus->name, __func__, drv->name, dev_name(dev));
25    563     if (!list_empty(&dev->devres_head)) {
26    564         dev_crit(dev, "Resources present before probing\n");
27    565         ret = -EBUSY;
28    566         goto done;
29    567     }
30    568
31    569 re_probe:
32    570     dev->driver = drv;
33    571
34    572     /* If using pinctrl, bind pins now before probing */
35    573     ret = pinctrl_bind_pins(dev);
36    574     if (ret)
37    575         goto pinctrl_bind_failed;
38    576
39    577     if (dev->bus->dma_configure) {
40    578         ret = dev->bus->dma_configure(dev);
41    579         if (ret)
42    580             goto probe_failed;
43    581     }
44    582
45    583     ret = driver_sysfs_add(dev);
46    584     if (ret) {
47    585         pr_err("%s: driver_sysfs_add(%s) failed\n",
48    586             __func__, dev_name(dev));
49    587         goto probe_failed;
50    588     }
51    589
52    590     if (dev->pm_domain && dev->pm_domain->activate) {
53    591         ret = dev->pm_domain->activate(dev);
54    592         if (ret)
55    593             goto probe_failed;
56    594     }

```

```

56 595
57 596     ret = call_driver_probe(dev, drv);
58 597     if (ret) {
59 598         /*
60 599          * Return probe errors as positive values so that the callers
61 600          * can distinguish them from other errors.
62 601          */
63 602         ret = -ret;
64 603         goto probe_failed;
65 604     }
66 605
67 606     ret = device_add_groups(dev, drv->dev_groups);
68 607     if (ret) {
69 608         dev_err(dev, "device_add_groups() failed\n");
70 609         goto dev_groups_failed;
71 610     }
72 611
73 612     if (dev_has_sync_state(dev)) {
74 613         ret = device_create_file(dev, &dev_attr_state_synced);
75 614         if (ret) {
76 615             dev_err(dev, "state_synced sysfs add failed\n");
77 616             goto dev_sysfs_state_synced_failed;
78 617         }
79 618     }
80 619
81 620     if (test_remove) {
82 621         test_remove = false;
83 622
84 623         device_remove_file(dev, &dev_attr_state_synced);
85 624         device_remove_groups(dev, drv->dev_groups);
86 625
87 626         if (dev->bus->remove)
88 627             dev->bus->remove(dev);
89 628         else if (drv->remove)
90 629             drv->remove(dev);
91 630
92 631         devres_release_all(dev);
93 632         driver_sysfs_remove(dev);
94 633         dev->driver = NULL;
95 634         dev_set_drvdata(dev, NULL);
96 635         if (dev->pm_domain && dev->pm_domain->dismiss)
97 636             dev->pm_domain->dismiss(dev);
98 637         pm_runtime_reinit(dev);
99 638
100 639         goto re_probe;
101 640     }
102 641
103 642     pinctrl_init_done(dev);
104 643
105 644     if (dev->pm_domain && dev->pm_domain->sync)
106 645         dev->pm_domain->sync(dev);
107 646
108 647     driver_bound(dev);
109 648     pr_debug("bus: '%s': %s: bound device %s to driver %s\n",
110 649             drv->bus->name, __func__, dev_name(dev), drv->name);
111 650     goto done;
112 651
113 652 dev_sysfs_state_synced_failed:
114 653     device_remove_groups(dev, drv->dev_groups);
115 654 dev_groups_failed:
116 655     if (dev->bus->remove)
117 656         dev->bus->remove(dev);
118 657     else if (drv->remove)
119 658         drv->remove(dev);
120 659 probe_failed:
121 660     if (dev->bus)
122 661         blocking_notifier_call_chain(&dev->bus->p->bus_notifier,
123 662                                     BUS_NOTIFY_DRIVER_NOT_BOUND, dev);
124 663 pinctrl_bind_failed:
125 664     device_links_no_driver(dev);
126 665     devres_release_all(dev);

```

```

127 | 666     arch_teardown_dma_ops(dev);
128 | 667     kfree(dev->dma_range_map);
129 | 668     dev->dma_range_map = NULL;
130 | 669     driver_sysfs_remove(dev);
131 | 670     dev->driver = NULL;
132 | 671     dev_set_drvdata(dev, NULL);
133 | 672     if (dev->pm_domain && dev->pm_domain->dismiss)
134 | 673         dev->pm_domain->dismiss(dev);
135 | 674     pm_runtime_reinit(dev);
136 | 675     dev_pm_set_driver_flags(dev, 0);
137 | 676 done:
138 | 677     return ret;
      | 678 }

```

1.3.14 call_driver_probe

code path: `drivers/base/dd.c`

```

1 | 510 static int call_driver_probe(struct device *dev, struct device_driver *drv)
2 | 511 {
3 | 512     int ret = 0;
4 | 513
5 | 514     if (dev->bus->probe)
6 | 515         ret = dev->bus->probe(dev);
7 | 516     else if (drv->probe)
8 | 517         ret = drv->probe(dev);
9 | 518
10 | 519     switch (ret) {
11 | 520     case 0:
12 | 521         break;
13 | 522     case -EPROBE_DEFER:
14 | 523         /* Driver requested deferred probing */
15 | 524         dev_dbg(dev, "Driver %s requests probe deferral\n", drv->name);
16 | 525         break;
17 | 526     case -ENODEV:
18 | 527     case -ENXIO:
19 | 528         pr_debug("%s: probe of %s rejects match %d\n",
20 | 529             drv->name, dev_name(dev), ret);
21 | 530         break;
22 | 531     default:
23 | 532         /* driver matched but the probe failed */
24 | 533         pr_warn("%s: probe of %s failed with error %d\n",
25 | 534             drv->name, dev_name(dev), ret);
26 | 535         break;
27 | 536     }
28 | 537
29 | 538     return ret;
30 | 539 }

```

1.3.15 platform_drv_probe

code path: `drivers/base/platform.c`

```

1 | 505 static int platform_drv_probe(struct device *_dev)
2 | 506 {
3 | 507     struct platform_driver *drv = to_platform_driver(_dev->driver);
4 | 508     struct platform_device *dev = to_platform_device(_dev);
5 | 509     int ret;
6 | 510
7 | 511     ret = of_clk_set_defaults(_dev->of_node, false);
8 | 512     if (ret < 0)
9 | 513         return ret;
10 | 514
11 | 515     ret = dev_pm_domain_attach(_dev, true);
12 | 516     if (ret != -EPROBE_DEFER) {
13 | 517         if (drv->probe) {
14 | 518             ret = drv->probe(dev); //point to rockchip_pcie_probe
15 |

```

```

16 |         519             if (ret)
17 |             520                 dev_pm_domain_detach(_dev, true);
18 |         521         } else {
19 |             522             /* don't fail if just dev_pm_domain_attach failed */
20 |             523             ret = 0;
21 |         524         }
22 |     525     }
23 |     526
24 |     527     if (drv->prevent_deferred_probe && ret == -EPROBE_DEFER) {
25 |         528         dev_warn(_dev, "probe deferral not supported\n");
26 |         529         ret = -ENXIO;
27 |     530     }
28 |     531
29 |     532     return ret;
    533 }

```

1.3.16 rockchip_pcie_probe

code path: `drivers/base/platform.c`

```

1 |         934 static int rockchip_pcie_probe(struct platform_device *pdev)
2 |         935 {
3 |             ...
4 |     1015 }

```

2. Device driver registration matching process:

In this process, the nvme driver is used as an example to introduce the matching process. nvme is a device mounted on the PCIe bus.

2.1 pci_register_driver

When pci_register_driver is registered, it will do the corresponding matching check according to the dev_driver and devices registered on the PCI bus. The whole process is as follows:

```

1 | |- nvme_init
2 |   |- pci_register_driver
3 |     |- __pci_register_driver
4 |       |- driver_register
5 |         |- bus_add_driver
6 |           |- driver_attach
7 |             |- bus_for_each_dev
8 |               |- __driver_attach
9 |                 |- driver_match_device
10 |                   |- pci_bus_match
11 |                     |- pci_match_device
12 |                       |- pci_match_one_device
13 |                         |- driver_probe_device
14 |                           |- really_probe
15 |                             |- call_driver_probe
16 |                               |- pci_device_probe
17 |                                 |- __pci_device_probe
18 |                                   |- pci_call_probe
19 |                                     |- local_pci_probe
20 |                                       |- nvme_probe

```

2.1.1 nvme_init

nvme device driver registration

code path: `drivers/nvme/host/pci.c`

```

1 |
2 |     3308 static int __init nvme_init(void)
3 |     3309 {
4 |         3310         BUILD_BUG_ON(sizeof(struct nvme_create_cq) != 64);
5 |         3311         BUILD_BUG_ON(sizeof(struct nvme_create_sq) != 64);
6 |         3312         BUILD_BUG_ON(sizeof(struct nvme_delete_queue) != 64);
7 |         3313         BUILD_BUG_ON(IRQ_AFFINITY_MAX_SETS < 2);

```

```

8      3314
9      3315     return pci_register_driver(&nvme_driver);
10     3316 }
11     3317
12     3318 static void __exit nvme_exit(void)
13     3319 {
14     3320     pci_unregister_driver(&nvme_driver);
15     3321     flush_workqueue(nvme_wq);
16     3322 }

```

2.1.2 pci_register_driver

code path: `include/linux/pci.h`

```

1      1400 /* pci_register_driver() must be a macro so KBUILD_MODNAME can be expanded */
2      1401 #define pci_register_driver(driver) \
3      1402     __pci_register_driver(driver, THIS_MODULE, KBUILD_MODNAME)

```

2.1.3 __pci_register_driver

`drv->driver.bus = &pci_bus_type;` indicates that the nvme device is a pci device mounted on the PCI bus, and the match function of `pci_bus_type` uses `feet` to detect whether the device mounted on the PCI/PCIe bus matches the driver.

code path: `drivers/pci/pci-driver.c`

```

1      1368 /**
2      1369  * __pci_register_driver - register a new pci driver
3      1370  * @drv: the driver structure to register
4      1371  * @owner: owner module of drv
5      1372  * @mod_name: module name string
6      1373  *
7      1374  * Adds the driver structure to the list of registered drivers.
8      1375  * Returns a negative value on error, otherwise 0.
9      1376  * If no error occurred, the driver remains registered even if
10     1377  * no device was claimed during registration.
11     1378  */
12     1379 int __pci_register_driver(struct pci_driver *drv, struct module *owner,
13     1380                          const char *mod_name)
14     1381 {
15     1382     /* initialize common driver fields */
16     1383     drv->driver.name = drv->name;
17     1384     drv->driver.bus = &pci_bus_type;
18     1385     drv->driver.owner = owner;
19     1386     drv->driver.mod_name = mod_name;
20     1387     drv->driver.groups = drv->groups;
21     1388     drv->driver.dev_groups = drv->dev_groups;
22     1389
23     1390     spin_lock_init(&drv->dynids.lock);
24     1391     INIT_LIST_HEAD(&drv->dynids.list);
25     1392
26     1393     /* register with core */
27     1394     return driver_register(&drv->driver);
28     1395 }
29     1396 EXPORT_SYMBOL(__pci_register_driver);

```

2.1.4 driver_register

code path: `drivers/base/driver.c`

```

1      139 /**
2      140  * driver_register - register driver with bus
3      141  * @drv: driver to register
4      142  *
5      143  * We pass off most of the work to the bus_add_driver() call,
6      144  * since most of the things we have to do deal with the bus
7      145  * structures.
8      146  */
9      147 int driver_register(struct device_driver *drv)
10     148 {
11     149     int ret;

```

```

11 150 struct device_driver *other;
12 151
13 152 if (!drv->bus->p) {
14 153     pr_err("Driver '%s' was unable to register with bus_type '%s' because the bus was not initialized.\n",
15 154         drv->name, drv->bus->name);
16 155     return -EINVAL;
17 156 }
18 157
19 158 if ((drv->bus->probe && drv->probe) ||
20 159     (drv->bus->remove && drv->remove) ||
twen 160     (drv->bus->shutdown && drv->shutdown))
twen 161     pr_warn("Driver '%s' needs updating - please use "
twen 162         "bus_type methods\n", drv->name);
twen 163
25 164 other = driver_find(drv->name, drv->bus);
26 165 if (other) {
27 166     pr_err("Error: Driver '%s' is already registered, "
28 167         "aborting...\n", drv->name);
29 168     return -EBUSY;
30 169 }
31 170
32 171 ret = bus_add_driver(drv);
33 172 if (ret)
34 173     return ret;
35 174 ret = driver_add_groups(drv, drv->groups);
36 175 if (ret) {
37 176     bus_remove_driver(drv);
38 177     return ret;
39 178 }
40 179 kobject_uevent(&drv->p->kobj, KOBJ_ADD);
41 180
42 181 return ret;
43 182 }
44 183 EXPORT_SYMBOL_GPL(driver_register);
45

```

2.1.5 bus_add_driver

code path: `drivers/base/bus.c`

```

1
2 586 /**
3 587  * bus_add_driver - Add a driver to the bus.
4 588  * @drv: driver.
5 589  */
6 590 int bus_add_driver(struct device_driver *drv)
7 591 {
8 592     struct bus_type *bus;
9 593     struct driver_private *priv;
10 594     int error = 0;
11 595
12 596     bus = bus_get(drv->bus);
13 597     if (!bus)
14 598         return -EINVAL;
15 599
16 600     pr_debug("bus: '%s': add driver %s\n", bus->name, drv->name);
17 601
18 602     priv = kzalloc(sizeof(*priv), GFP_KERNEL);
19 603     if (!priv) {
20 604         error = -ENOMEM;
twen 605         goto out_put_bus;
twen 606     }
twen 607     klist_init(&priv->klist_devices, NULL, NULL);
twen 608     priv->driver = drv;
twen 609     drv->p = priv;
25 610     priv->kobj.kset = bus->p->drivers_kset;
26 611     error = kobject_init_and_add(&priv->kobj, &driver_ktype, NULL,
27 612         "%s", drv->name);
28 613     if (error)
29 614         goto out_unregister;
30 615
31

```

```

32     616     klist_add_tail(&priv->knode_bus, &bus->p->klist_drivers);
33     617     if (drv->bus->p->drivers_autoprobe) {
34     618         error = driver_attach(drv);
35     619         if (error)
36     620             goto out_unregister;
37     621     }
38     622     module_add_driver(drv->owner, drv);
39     623
40     624     error = driver_create_file(drv, &driver_attr_uevent);
41     625     if (error) {
42     626         printk(KERN_ERR "%s: uevent attr (%s) failed\n",
43     627             __func__, drv->name);
44     628     }
45     629     error = driver_add_groups(drv, bus->drv_groups);
46     630     if (error) {
47     631         /* How the hell do we get out of this pickle? Give up */
48     632         printk(KERN_ERR "%s: driver_add_groups(%s) failed\n",
49     633             __func__, drv->name);
50     634     }
51     635
52     636     if (!drv->suppress_bind_attrs) {
53     637         error = add_bind_files(drv);
54     638         if (error) {
55     639             /* Ditto */
56     640             printk(KERN_ERR "%s: add_bind_files(%s) failed\n",
57     641                 __func__, drv->name);
58     642         }
59     643     }
60     644
61     645     return 0;
62     646
63     647 out_unregister:
64     648     kobject_put(&priv->kobj);
65     649     /* drv->p is freed in driver_release() */
66     650     drv->p = NULL;
67     651 out_put_bus:
68     652     bus_put(bus);
69     653     return error;
654 }

```

2.1.6 driver_attach

code path: `drivers/base/dd.c`

```

1     1146 /**
2     1147  * driver_attach - try to bind driver to devices.
3     1148  * @drv: driver.
4     1149  *
5     1150  * Walk the list of devices that the bus has on it and try to
6     1151  * match the driver with each one. If driver_probe_device()
7     1152  * returns 0 and the @dev->driver is set, we've found a
8     1153  * compatible pair.
9     1154  */
10    1155 int driver_attach(struct device_driver *drv)
11    1156 {
12    1157     return bus_for_each_dev(drv->bus, NULL, drv, __driver_attach);
13    1158 }
14    1159 EXPORT_SYMBOL_GPL(driver_attach);

```

2.1.7 bus_for_each_dev

code path: `drivers/base/bus.c`

```

1     269 /**
2     270  * bus_for_each_dev - device iterator.
3     271  * @bus: bus type.
4     272  * @start: device to start iterating from.
5     273  * @data: data for the callback.
6     274  * @fn: function to be called for each device.
7     275  *
8     276  * Iterate over @bus's list of devices, and call @fn for each,

```



```

8      277 * passing it @data. If @start is not NULL, we use that device to
9      278 * begin iterating from.
10     279 *
11     280 * We check the return of @fn each time. If it returns anything
12     281 * other than 0, we break out and return that value.
13     282 *
14     283 * NOTE: The device that returns a non-zero value is not retained
15     284 * in any way, nor is its refcount incremented. If the caller needs
16     285 * to retain this data, it should do so, and increment the reference
17     286 * count in the supplied callback.
18     287 */
19     288 int bus_for_each_dev(struct bus_type *bus, struct device *start,
20     289                     void *data, int (*fn)(struct device *, void *))
21     290 {
22     291     struct klist_iter i;
23     292     struct device *dev;
24     293     int error = 0;
25     294
26     295     if (!bus || !bus->p)
27     296         return -EINVAL;
28     297
29     298     klist_iter_init_node(&bus->p->klist_devices, &i,
30     299                         (start ? &start->p->knode_bus : NULL));
31     300     while (!error && (dev = next_device(&i)))
32     301         error = fn(dev, data);
33     302     klist_iter_exit(&i);
34     303     return error;
35     304 }
36     305 EXPORT_SYMBOL_GPL(bus_for_each_dev);

```

2.1.8 __driver_attach

code path: [drivers/base/dd.c](#)

```

1      1092 static int __driver_attach(struct device *dev, void *data)
2      1093 {
3      1094     struct device_driver *drv = data;
4      1095     int ret;
5      1096
6      1097     /*
7      1098      * Lock device and try to bind to it. We drop the error
8      1099      * here and always return 0, because we need to keep trying
9      1100      * to bind to devices and some drivers will return an error
10     1101      * simply if it didn't support the device.
11     1102      *
12     1103      * driver_probe_device() will spit a warning if there
13     1104      * is an error.
14     1105      */
15     1106
16     1107     ret = driver_match_device(drv, dev);
17     1108     if (ret == 0) {
18     1109         /* no match */
19     1110         return 0;
20     1111     } else if (ret == -EPROBE_DEFER) {
21     1112         dev_dbg(dev, "Device match requests probe deferral\n");
22     1113         dev->can_match = true;
23     1114         driver_deferred_probe_add(dev);
24     1115     } else if (ret < 0) {
25     1116         dev_dbg(dev, "Bus failed to match device: %d\n", ret);
26     1117         return ret;
27     1118     } /* ret > 0 means positive match */
28     1119
29     1120     if (driver_allows_async_probing(drv)) {
30     1121         /*
31     1122          * Instead of probing the device synchronously we will
32     1123          * probe it asynchronously to allow for more parallelism.
33     1124          *
34     1125          * We only take the device lock here in order to guarantee
35     1126          * that the dev->driver and async_driver fields are protected
36     1127          */
37     1128         dev_dbg(dev, "probing driver %s asynchronously\n", drv->name);
38     1129         device_lock(dev);
39     1130         if (!dev->driver) {
40     1131             get_device(dev);
41     1132             dev->p->async_driver = drv;

```

```

40      1133         async_schedule_dev(__driver_attach_async_helper, dev);
41      1134     }
42      1135     device_unlock(dev);
43      1136     return 0;
44      1137 }
45      1138
46      1139     __device_driver_lock(dev, dev->parent);
47      1140     driver_probe_device(drv, dev);
48      1141     __device_driver_unlock(dev, dev->parent);
49      1142
50      1143     return 0;
51      1144 }
52
53

```

2.1.9 driver_match_device

The match function of `drv->bus->match(dev, drv)` points to the `pci_bus_match` function defined by the previously registered `pci_bus_type`.

code path: `drivers/base/base.h`

```

1      144 static inline int driver_match_device(struct device_driver *drv,
2      145         struct device *dev)
3      146 {
4      147     return drv->bus->match ? drv->bus->match(dev, drv) : 1;
5      148

```

2.1.10 pci_bus_match

code path: `drivers/pci/pci-driver.c`

```

1      1440 /**
2      1441  * pci_bus_match - Tell if a PCI device structure has a matching PCI device id structure
3      1442  * @dev: the PCI device structure to match against
4      1443  * @drv: the device driver to search for matching PCI device id structures
5      1444  *
6      1445  * Used by a driver to check whether a PCI device present in the
7      1446  * system is in its list of supported devices. Returns the matching
8      1447  * pci_device_id structure or %NULL if there is no match.
9      1448  */
10     1449 static int pci_bus_match(struct device *dev, struct device_driver *drv)
11     1450 {
12     1451     struct pci_dev *pci_dev = to_pci_dev(dev);
13     1452     struct pci_driver *pci_drv;
14     1453     const struct pci_device_id *found_id;
15     1454
16     1455     if (!pci_dev->match_driver)
17     1456         return 0;
18     1457
19     1458     pci_drv = to_pci_driver(drv);
20     1459     found_id = pci_match_device(pci_drv, pci_dev);
21     1460     if (found_id)
22     1461         return 1;
23     1462
24     1463     return 0;
25     1464 }

```

2.1.11 pci_match_device

code path: `drivers/pci/pci-driver.c`

```

1      125 /**
2      126  * pci_match_device - See if a device matches a driver's list of IDs
3      127  * @drv: the PCI driver to match against
4      128  * @dev: the PCI device structure to match against
5      129  *
6      130  * Used by a driver to check whether a PCI device is in its list of
7      131  * supported devices or in the dynids list, which may have been augmented
8      132  * via the sysfs "new_id" file. Returns the matching pci_device_id
9      133  * structure or %NULL if there is no match.
10     134  */
11     135 static const struct pci_device_id *pci_match_device(struct pci_driver *drv,

```

```

11      136          struct pci_dev *dev)
12      137 {
13      138      struct pci_dynid *dynid;
14      139      const struct pci_device_id *found_id = NULL;
15      140
16      141      /* When driver_override is set, only bind to the matching driver */
17      142      if (dev->driver_override && strcmp(dev->driver_override, drv->name))
18      143          return NULL;
19      144
20      145      /* Look at the dynamic ids first, before the static ones */
twen      146      spin_lock(&drv->dynids.lock);
twen      147      list_for_each_entry(dynid, &drv->dynids.list, node) {
twen      148          if (pci_match_one_device(&dynid->id, dev)) {
twen      149              found_id = &dynid->id;
25      150              break;
26      151          }
27      152      }
28      153      spin_unlock(&drv->dynids.lock);
29      154
30      155      if (!found_id)
31      156          found_id = pci_match_id(drv->id_table, dev);
32      157
33      158      /* driver_override will always match, send a dummy id */
34      159      if (!found_id && dev->driver_override)
35      160          found_id = &pci_device_id_any;
36      161
37      162      return found_id;
38      163 }
39

```

2.1.12 pci_match_one_device

code path: [drivers/pci/pci-driver.c](#)

```

1      204 /**
2      205  * pci_match_one_device - Tell if a PCI device structure has a matching
3      206  *                        PCI device id structure
4      207  * @id: single PCI device id structure to match
5      208  * @dev: the PCI device structure to match against
6      209  *
7      210  * Returns the matching pci_device_id structure or %NULL if there is no match.
8      211  */
9      212 static inline const struct pci_device_id *
10     213 pci_match_one_device(const struct pci_device_id *id, const struct pci_dev *dev)
11     214 {
12     x 215     if ((id->vendor == PCI_ANY_ID || id->vendor == dev->vendor) &&
13     x 216         (id->device == PCI_ANY_ID || id->device == dev->device) &&
14     x 217         (id->subvendor == PCI_ANY_ID || id->subvendor == dev->subsystem_vendor) &&
15     x 218         (id->subdevice == PCI_ANY_ID || id->subdevice == dev->subsystem_device) &&
16     x 219         !((id->class ^ dev->class) & id->class_mask))
17     x 220         return id;
18     x 221     return NULL;
19     222 }

```

2.1.13 driver_probe_device

code path: [drivers/base/dd.c](#)

```

1
2      761 /**
3      762  * driver_probe_device - attempt to bind device & driver together
4      763  * @drv: driver to bind a device to
5      764  * @dev: device to try to bind to the driver
6      765  *
7      766  * This function returns -ENODEV if the device is not registered, -EBUSY if it
8      767  * already has a driver, 0 if the device is bound successfully and a positive
9      768  * (inverted) error code for failures from the ->probe method.
10     769  *
11     770  * This function must be called with @dev lock held. When called for a
12     771  * USB interface, @dev->parent lock must be held as well.
13     772  *
14     773  * If the device has a parent, runtime-resume the parent before driver probing.
15     774  */
16     775 static int driver_probe_device(struct device_driver *drv, struct device *dev)

```

```

14 776 {
15 777     int trigger_count = atomic_read(&deferred_trigger_count);
16 778     int ret;
17 779
18 780     atomic_inc(&probe_count);
19 781     ret = __driver_probe_device(drv, dev);
20 782     if (ret == -EPROBE_DEFER || ret == EPROBE_DEFER) {
21 783         driver_deferred_probe_add(dev);
22 784
23 785         /*
24 786          * Did a trigger occur while probing? Need to re-trigger if yes
25 787          */
26 788         if (trigger_count != atomic_read(&deferred_trigger_count) &&
27 789             !defer_all_probes)
28 790             driver_deferred_probe_trigger();
29 791     }
30 792     atomic_dec(&probe_count);
31 793     wake_up_all(&probe_waitqueue);
32 794     return ret;
33 795 }
34

```

2.1.12 __driver_probe_device

code path: `drivers/base/dd.c`

```

1 730 static int __driver_probe_device(struct device_driver *drv, struct device *dev)
2 731 {
3 732     int ret = 0;
4 733
5 734     if (dev->p->dead || !device_is_registered(dev))
6 735         return -ENODEV;
7 736     if (dev->driver)
8 737         return -EBUSY;
9 738
10 739     dev->can_match = true;
11 740     pr_debug("bus: '%s': %s: matched device %s with driver %s\n",
12 741         drv->bus->name, __func__, dev_name(dev), drv->name);
13 742
14 743     pm_runtime_get_suppliers(dev);
15 744     if (dev->parent)
16 745         pm_runtime_get_sync(dev->parent);
17 746
18 747     pm_runtime_barrier(dev);
19 748     if (initcall_debug)
20 749         ret = really_probe_debug(dev, drv);
21 750     else
22 751         ret = really_probe(dev, drv);
23 752     pm_request_idle(dev);
24 753
25 754     if (dev->parent)
26 755         pm_runtime_put(dev->parent);
27 756
28 757     pm_runtime_put_suppliers(dev);
29 758     return ret;
30 759 }

```

2.1.14 really_probe

code path: `drivers/base/dd.c`

```

2 541 static int really_probe(struct device *dev, struct device_driver *drv)
3 542 {
4 543     bool test_remove = IS_ENABLED(CONFIG_DEBUG_TEST_DRIVER_REMOVE) &&
5 544         !drv->suppress_bind_attrs;
6 545     int ret;
7 546
8 547     if (defer_all_probes) {

```

```

10      548      /*
11      549      * Value of defer_all_probes can be set only by
12      550      * device_block_probing() which, in turn, will call
13      551      * wait_for_device_probe() right after that to avoid any races.
14      552      */
15      553      dev_dbg(dev, "Driver %s force probe deferral\n", drv->name);
16      554      return -EPROBE_DEFER;
17      555  }
18      556
19      557  ret = device_links_check_suppliers(dev);
20      558  if (ret)
21      559      return ret;
22      560
23      561  pr_debug("bus: '%s': %s: probing driver %s with device %s\n",
24      562          drv->bus->name, __func__, drv->name, dev_name(dev));
25      563  if (!list_empty(&dev->devres_head)) {
26      564      dev_crit(dev, "Resources present before probing\n");
27      565      ret = -EBUSY;
28      566      goto done;
29      567  }
30      568
31      569  re_probe:
32      570      dev->driver = drv;
33      571
34      572      /* If using pinctrl, bind pins now before probing */
35      573      ret = pinctrl_bind_pins(dev);
36      574      if (ret)
37      575          goto pinctrl_bind_failed;
38      576
39      577      if (dev->bus->dma_configure) {
40      578          ret = dev->bus->dma_configure(dev);
41      579          if (ret)
42      580              goto probe_failed;
43      581      }
44      582
45      583      ret = driver_sysfs_add(dev);
46      584      if (ret) {
47      585          pr_err("%s: driver_sysfs_add(%s) failed\n",
48      586              __func__, dev_name(dev));
49      587          goto probe_failed;
50      588      }
51      589
52      590      if (dev->pm_domain && dev->pm_domain->activate) {
53      591          ret = dev->pm_domain->activate(dev);
54      592          if (ret)
55      593              goto probe_failed;
56      594      }
57      595
58      596      ret = call_driver_probe(dev, drv);
59      597      if (ret) {
60      598          /*
61      599          * Return probe errors as positive values so that the callers
62      600          * can distinguish them from other errors.
63      601          */
64      602          ret = -ret;
65      603          goto probe_failed;
66      604      }
67      605
68      606      ret = device_add_groups(dev, drv->dev_groups);
69      607      if (ret) {
70      608          dev_err(dev, "device_add_groups() failed\n");
71      609          goto dev_groups_failed;
72      610      }
73      611
74      612      if (dev_has_sync_state(dev)) {
75      613          ret = device_create_file(dev, &dev_attr_state_synced);
76      614          if (ret) {
77      615              dev_err(dev, "state_synced sysfs add failed\n");
78      616              goto dev_sysfs_state_synced_failed;
79      617          }
80      618      }
81      619

```

```

81      620     if (test_remove) {
82      621         test_remove = false;
83      622
84      623         device_remove_file(dev, &dev_attr_state_synced);
85      624         device_remove_groups(dev, drv->dev_groups);
86      625
87      626         if (dev->bus->remove)
88      627             dev->bus->remove(dev);
89      628         else if (drv->remove)
90      629             drv->remove(dev);
91      630
92      631         devres_release_all(dev);
93      632         driver_sysfs_remove(dev);
94      633         dev->driver = NULL;
95      634         dev_set_drvdata(dev, NULL);
96      635         if (dev->pm_domain && dev->pm_domain->dismiss)
97      636             dev->pm_domain->dismiss(dev);
98      637         pm_runtime_reinit(dev);
99      638
100     639         goto re_probe;
101     640     }
102     641
103     642     pinctrl_init_done(dev);
104     643
105     644     if (dev->pm_domain && dev->pm_domain->sync)
106     645         dev->pm_domain->sync(dev);
107     646
108     647     driver_bound(dev);
109     648     pr_debug("bus: '%s': %s: bound device %s to driver %s\n",
110     649         drv->bus->name, __func__, dev_name(dev), drv->name);
111     650     goto done;
112     651
113     652 dev_sysfs_state_synced_failed:
114     653     device_remove_groups(dev, drv->dev_groups);
115     654 dev_groups_failed:
116     655     if (dev->bus->remove)
117     656         dev->bus->remove(dev);
118     657     else if (drv->remove)
119     658         drv->remove(dev);
120     659 probe_failed:
121     660     if (dev->bus)
122     661         blocking_notifier_call_chain(&dev->bus->p->bus_notifier,
123     662             BUS_NOTIFY_DRIVER_NOT_BOUND, dev);
124     663 pinctrl_bind_failed:
125     664     device_links_no_driver(dev);
126     665     devres_release_all(dev);
127     666     arch_teardown_dma_ops(dev);
128     667     kfree(dev->dma_range_map);
129     668     dev->dma_range_map = NULL;
130     669     driver_sysfs_remove(dev);
131     670     dev->driver = NULL;
132     671     dev_set_drvdata(dev, NULL);
133     672     if (dev->pm_domain && dev->pm_domain->dismiss)
134     673         dev->pm_domain->dismiss(dev);
135     674     pm_runtime_reinit(dev);
136     675     dev_pm_set_driver_flags(dev, 0);
137     676 done:
138     677     return ret;
139     678 }

```

2.1.15 call_driver_probe

For devices mounted on the PCI/PCIe bus, when call_driver_probe is called, the probe function of the corresponding device driver will be called through `ret = drv->probe(dev)`.

code path: `drivers/base/dd.c`

```

1
2     510 static int call_driver_probe(struct device *dev, struct device_driver *drv)
3     511 {

```

```

4      512     int ret = 0;
5      513
6      514     if (dev->bus->probe)
7      515         ret = dev->bus->probe(dev);
8      516     else if (drv->probe)
9      517         ret = drv->probe(dev);
10     518
11     519     switch (ret) {
12     520     case 0:
13     521         break;
14     522     case -EPROBE_DEFER:
15     523         /* Driver requested deferred probing */
16     524         dev_dbg(dev, "Driver %s requests probe deferral\n", drv->name);
17     525         break;
18     526     case -ENODEV:
19     527     case -ENXIO:
20     528         pr_debug("%s: probe of %s rejects match %d\n",
twenty 529             drv->name, dev_name(dev), ret);
twenty 530         break;
twenty 531     default:
twenty 532         /* driver matched but the probe failed */
25     533         pr_warn("%s: probe of %s failed with error %d\n",
26     534             drv->name, dev_name(dev), ret);
27     535         break;
28     536     }
29     537
30     538     return ret;
        539 }

```

2.1.16 pci_bus_type

code path: `drivers/pci/pci-driver.c`

```

1      1600 struct bus_type pci_bus_type = {
2      1601     .name      = "pci",
3      1602     .match      = pci_bus_match,
4      1603     .uevent      = pci_uevent,
5      1604     .probe      = pci_device_probe,
6      1605     .remove      = pci_device_remove,
7      1606     .shutdown    = pci_device_shutdown,
8      1607     .dev_groups  = pci_dev_groups,
9      1608     .bus_groups  = pci_bus_groups,
10     1609     .drv_groups  = pci_drv_groups,
11     1610     .pm          = PCI_PM_OPS_PTR,
12     1611     .num_vf      = pci_bus_num_vf,
13     1612     .dma_configure = pci_dma_configure,
14     1613 };
15     1614 EXPORT_SYMBOL(pci_bus_type);

```

2.1.17 pci_device_probe

`pci_device_probe` corresponds to `drv->probe(dev)`, and the probe function of the corresponding device driver will be called through the `pci_device_probe` function.

code path: `drivers/pci/pci-driver.c`

```

1      418 static int pci_device_probe(struct device *dev)
2      419 {
3      420     int error;
4      421     struct pci_dev *pci_dev = to_pci_dev(dev);
5      422     struct pci_driver *drv = to_pci_driver(dev->driver);
6      423
7      424     if (!pci_device_can_probe(pci_dev))
8      425         return -ENODEV;
9      426
10     427     pci_assign_irq(pci_dev);
11     428
12     429     error = pcibios_alloc_irq(pci_dev);

```

```

14      430     if (error < 0)
15      431         return error;
16      432
17      433     pci_dev_get(pci_dev);
18      434     error = __pci_device_probe(drv, pci_dev);
19      435     if (error) {
20      436         pcibios_free_irq(pci_dev);
twen    437         pci_dev_put(pci_dev);
twen    438     }
twen    439
twen    440     return error;
twen    441 }

```

2.1.18 __pci_device_probe

code path: `drivers/pci/pci-driver.c`

```

1      373 /**
2      374  * __pci_device_probe - check if a driver wants to claim a specific PCI device
3      375  * @drv: driver to call to check if it wants the PCI device
4      376  * @pci_dev: PCI device being probed
5      377  *
6      378  * returns 0 on success, else error.
7      379  * side-effect: pci_dev->driver is set to drv when drv claims pci_dev.
8      380  */
9      381 static int __pci_device_probe(struct pci_driver *drv, struct pci_dev *pci_dev)
10     382 {
11     383     const struct pci_device_id *id;
12     384     int error = 0;
13     385
14     386     if (!pci_dev->driver && drv->probe) {
15     387         error = -ENODEV;
16     388
17     389         id = pci_match_device(drv, pci_dev); -->同2.1.11的代码流程分析
18     390         if (id)
19     391             error = pci_call_probe(drv, pci_dev, id);
20     392     }
twen    393     return error;
twen    394 }

```

2.1.19 pci_call_probe

code path: `drivers/pci/pci-driver.c`

```

1      335 static int pci_call_probe(struct pci_driver *drv, struct pci_dev *dev,
2      336                          const struct pci_device_id *id)
3      337 {
4      338     int error, node, cpu;
5      339     int hk_flags = HK_FLAG_DOMAIN | HK_FLAG_WQ;
6      340     struct drv_dev_and_id ddi = { drv, dev, id };
7      341
8      342     /*
9      343      * Execute driver initialization on node where the device is
10     344      * attached. This way the driver likely allocates its local memory
11     345      * on the right node.
12     346      */
13     347     node = dev_to_node(&dev->dev);
14     348     dev->is_probed = 1;
15     349
16     350     cpu_hotplug_disable();
17     351
18     352     /*
19     353      * Prevent nesting work_on_cpu() for the case where a Virtual Function
20     354      * device is probed from work_on_cpu() of the Physical device.
twen    355      */
twen    356     if (node < 0 || node >= MAX_NUMNODES || !node_online(node) ||
twen    357         pci_physfn_is_probed(dev))
twen    358         cpu = nr_cpu_ids;
25     359     else
26     360         cpu = cpumask_any_and(cpumask_of_node(node),

```



```

27 | 361             housekeeping_cpumask(hk_flags));
28 | 362
29 | 363     if (cpu < nr_cpu_ids)
30 | 364         error = work_on_cpu(cpu, local_pci_probe, &ddi);
31 | 365     else
32 | 366         error = local_pci_probe(&ddi);
33 | 367
34 | 368     dev->is_probed = 0;
35 | 369     cpu_hotplug_enable();
36 | 370     return error;
37 | 371 }

```

2.1.20 local_pci_probe

code path: `drivers/pci/pci-driver.c`

```

1 | 290 static long local_pci_probe(void *_ddi)
2 | 291 {
3 | 292     struct drv_dev_and_id *ddi = _ddi;
4 | 293     struct pci_dev *pci_dev = ddi->dev;
5 | 294     struct pci_driver *pci_drv = ddi->drv;
6 | 295     struct device *dev = &pci_dev->dev;
7 | 296     int rc;
8 | 297
9 | 298     /*
10 | 299      * Unbound PCI devices are always put in D0, regardless of
11 | 300      * runtime PM status. During probe, the device is set to
12 | 301      * active and the usage count is incremented. If the driver
13 | 302      * supports runtime PM, it should call pm_runtime_put_noidle(),
14 | 303      * or any other runtime PM helper function decrementing the usage
15 | 304      * count, in its probe routine and pm_runtime_get_noresume() in
16 | 305      * its remove routine.
17 | 306      */
18 | 307     pm_runtime_get_sync(dev);
19 | 308     pci_dev->driver = pci_drv;
20 | 309     rc = pci_drv->probe(pci_dev, ddi->id);
21 | 310     if (!rc)
22 | 311         return rc;
23 | 312     if (rc < 0) {
24 | 313         pci_dev->driver = NULL;
25 | 314         pm_runtime_put_sync(dev);
26 | 315         return rc;
27 | 316     }
28 | 317     /*
29 | 318      * Probe function should return < 0 for failure, 0 for success
30 | 319      * Treat values > 0 as success, but warn.
31 | 320      */
32 | 321     pci_warn(pci_dev, "Driver probe function unexpectedly returned %d\n",
33 | 322              rc);
34 | 323     return 0;
35 | 324 }

```

2.1.21 nvme_probe

```

1 | 2882 static int nvme_probe(struct pci_dev *pdev, const struct pci_device_id *id)
2 | 2883 {
3 | 2884     ...
4 | 2885 }

```

The knowledge points of the article are matched with the official knowledge files, and relevant knowledge can be further learned